

Final Audit Report

Executive Summary

Auto-generated audit report by Chief Justice.

Architecture Deep Dive

(No architecture deep-dive found in provided PDF; please add sections on Dialectical Synthesis, Fan-In/Fan-Out, and Metacognition)

Architectural Diagrams

Criterion-by-Criterion Breakdown

Theoretical Depth (Documentation) â Final Score: 2

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Cited Evidence: Determine theoretical depth
- Prosecutor (1): Missing evidence: Determine theoretical depth
- Cited Evidence: Determine theoretical depth
- TechLead (2): No clear artifacts; technical debt suspected.
- Cited Evidence: Determine theoretical depth

Remediation: See detective evidence and implement missing artifacts.

Report Accuracy (Cross-Reference) â Final Score: 2

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Cited Evidence: Extract file paths from PDF
- Prosecutor (1): Missing evidence: Extract file paths from PDF
- Cited Evidence: Extract file paths from PDF
- TechLead (2): No clear artifacts; technical debt suspected.
- Cited Evidence: Extract file paths from PDF

Remediation: See detective evidence and implement missing artifacts.

Git Forensic Analysis â Final Score: 4

Dissent: High variance across judges; median score used after re-evaluation.

Judge Opinions:

- Defense (5): Evidence of intent and partial implementation found; reward effort and intent.

- Cited Evidence: Extract Git History Progression
- Prosecutor (2): Found concerning patterns or insufficient evidence.
- Cited Evidence: Extract Git History Progression
- TechLead (4): Artifacts present; pragmatic functionality likely.
- Cited Evidence: Extract Git History Progression

Remediation: See detective evidence and implement missing artifacts.

State Management Rigor â Final Score: 5

Dissent: High variance across judges; median score used after re-evaluation.

Judge Opinions:

- Defense (5): Evidence of intent and partial implementation found; reward effort and intent.
- Cited Evidence: Verify State File Existence, Verify Pydantic/TypedDict usage, Verify Reducers
- Prosecutor (2): Found concerning patterns or insufficient evidence.
- Cited Evidence: Verify State File Existence, Verify Pydantic/TypedDict usage, Verify Reducers
- TechLead (5): Proper reducers detected; good parallel safety.
- Cited Evidence: Verify State File Existence, Verify Pydantic/TypedDict usage, Verify Reducers

Remediation: Define AgentState with Pydantic or TypedDict and use Annotated reducers (operator.add, operator.ior) to avoid p

Graph Orchestration Architecture â Final Score: 4

Dissent: High variance across judges; median score used after re-evaluation.

Judge Opinions:

- Defense (5): Evidence of intent and partial implementation found; reward effort and intent.
- Cited Evidence: Verify StateGraph Definition, Verify Fan-Out / Fan-In patterns
- Prosecutor (2): Found concerning patterns or insufficient evidence.
- Cited Evidence: Verify StateGraph Definition, Verify Fan-Out / Fan-In patterns
- TechLead (4): Artifacts present; pragmatic functionality likely.
- Cited Evidence: Verify StateGraph Definition, Verify Fan-Out / Fan-In patterns

Remediation: Implement parallel fan-out for Detectives and Judges with a fan-in EvidenceAggregator node; add conditional edg

Safe Tool Engineering â Final Score: 1

Dissent: High variance and re-evaluation found missing cited evidence; prosecutor position favored.

Judge Opinions:

- Defense (5): Evidence of intent and partial implementation found; reward effort and intent.
- Cited Evidence: Verify Git Sandboxing, Security Violations
- Prosecutor (1): Missing evidence: Security Violations; Raw os.system usage detected: security risk
- Cited Evidence: Verify Git Sandboxing, Security Violations
- TechLead (5): Sandboxed cloning detected.
- Cited Evidence: Verify Git Sandboxing, Security Violations

Remediation: Ensure git clone uses tempfile.TemporaryDirectory and subprocess.run with error handling; remove raw os.system

Structured Output Enforcement â Final Score: 2

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Cited Evidence: Structured Output Usage
- Prosecutor (1): Missing evidence: Structured Output Usage
- Cited Evidence: Structured Output Usage
- TechLead (2): No clear artifacts; technical debt suspected.
- Cited Evidence: Structured Output Usage

Remediation: See detective evidence and implement missing artifacts.

Judicial Nuance and Dialectics â Final Score: 3

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Prosecutor (3): Found concerning patterns or insufficient evidence.
- TechLead (2): No clear artifacts; technical debt suspected.

Remediation: See detective evidence and implement missing artifacts.

Chief Justice Synthesis Engine â Final Score: 3

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Prosecutor (3): Found concerning patterns or insufficient evidence.
- TechLead (2): No clear artifacts; technical debt suspected.

Remediation: See detective evidence and implement missing artifacts.

Architectural Diagram Analysis â Final Score: 2

Judge Opinions:

- Defense (3): No direct evidence found in this dimension, but allow mitigation for effort shown elsewhere.
- Cited Evidence: Architectural Diagram Analysis
- Prosecutor (1): Missing evidence: Architectural Diagram Analysis
- Cited Evidence: Architectural Diagram Analysis
- TechLead (2): No clear artifacts; technical debt suspected.
- Cited Evidence: Architectural Diagram Analysis

Remediation: See detective evidence and implement missing artifacts.

Reflection on the MinMax Feedback Loop

(Add reflections: what peer agents caught, how you updated your agent, and remaining gaps.)

Remediation Plan

See per-criterion remediation above.